

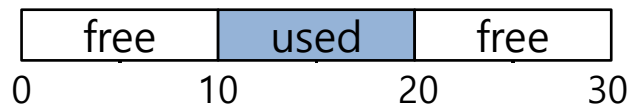
# 17. Free-Space Management

Operating System: Three Easy Pieces

---

# Variable sized allocation

- Free space management is easier when the memory is divided into fixed size units
- For variable-sized units, the problem that exists is known as **external fragmentation**
  - Arises in a user-level memory-allocation library (as in `malloc()` and `free()`)
  - And in an OS managing physical memory using segmentation
  - A request of 15 bytes will fail here even though total 20 bytes is free




Given a block of memory, how do we allocate variable sized memory allocation requests to minimize fragmentation and overheads ?


# Common Low Level Mechanism

- ▣ **Free List:** generic data structure used by the library to manage free space in the heap

- ◆ Assume the following 30 byte heap



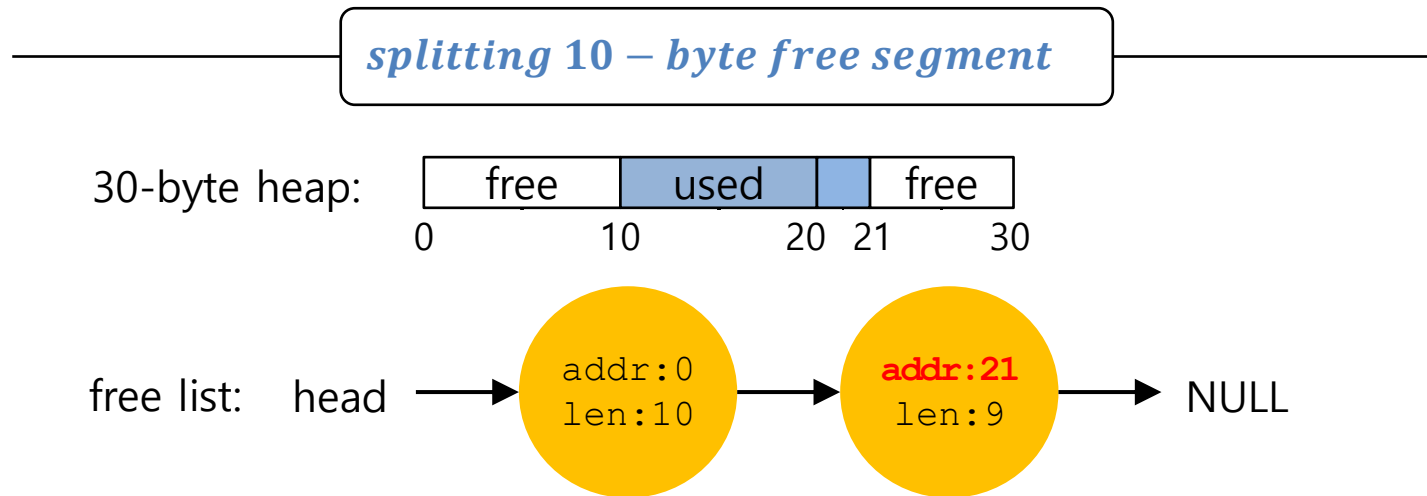
- ◆ The free list for this could be



- ◆ Any request = 10 bytes could be satisfied by either
- ◆ What happens if the request < 10 bytes ?


# Common Low Level Mechanism

- ▣ **Splitting**: Finding a free chunk of memory that can satisfy the request and splitting it into two.
  - ◆ In previous example if we have **1-byte request** and let's say the allocator decided to use the second of the two elements on the list
  - ◆ That chunk would be split into two, returning the first chunk of 1-byte allocated region (with the address 20) and the second chunk would remain in list.



# Common Low Level Mechanism - Coalescing

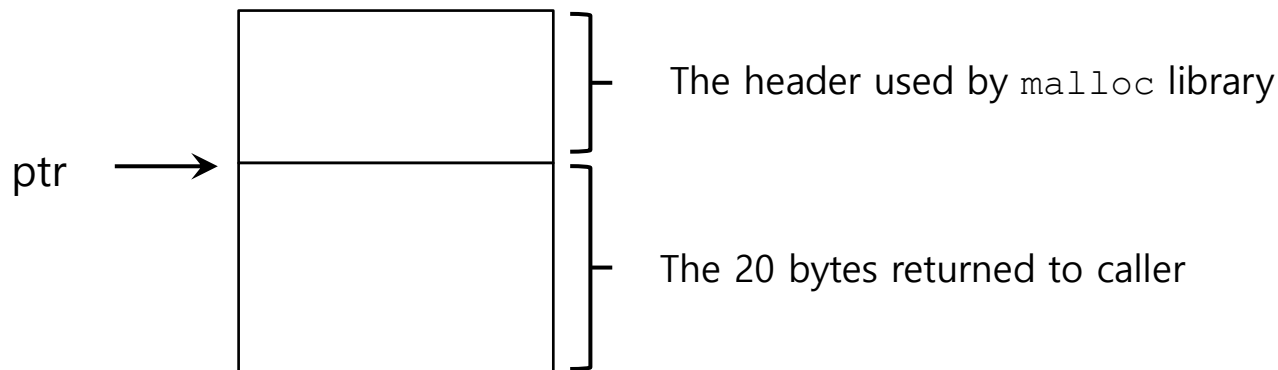
- ❑ If a user requests memory that is **bigger than free chunk size**, the list will **not find** such a free chunk.
- ❑ **Coalescing: Merge** existing free chunks into a large single free chunk if **addresses** of them are **nearby**. Ex. If the application calls `free(10)` returning the middle space



# Tracking The Size of Allocated Regions


- ▣ The interface to `free(void *ptr)` does **not take a size parameter**.
  - ◆ How does the library **know the size** of memory region that will be back **into free list**?
- ▣ Most allocators store **extra information** in a **header block**.

```
ptr = malloc(20);
```



# The Header of Allocated Memory Chunk

- ▣ The header minimally **contains the size** of the allocated memory region.
- ▣ The header may also contain
  - ◆ A magic number for integrity checking
  - ◆ Additional pointers to speed up deallocation



```
typedef struct __header_t {  
    int size;  
    int magic;  
} header_t;
```

A Simple Header

## The Header of Allocated Memory Chunk(Cont.)

- ▣ If a user **request  $N$  bytes**, the library searches for a free chunk of **size  $N$  plus the size of the header**
- ▣ When user calls **free (ptr)**, library uses simple pointer arithmetic to find the header pointer.


```
void free(void *ptr) {  
    header_t *hptr = (void *)ptr - sizeof(header_t);  
}
```



# Embedding a Free List


- Free space managed as a linked list
  - Pointer to the next free chunk is embedded within the free chunk
- The library tracks the head of the list
  - Allocations happen from the head

A 4KB Heap With One Free Chunk



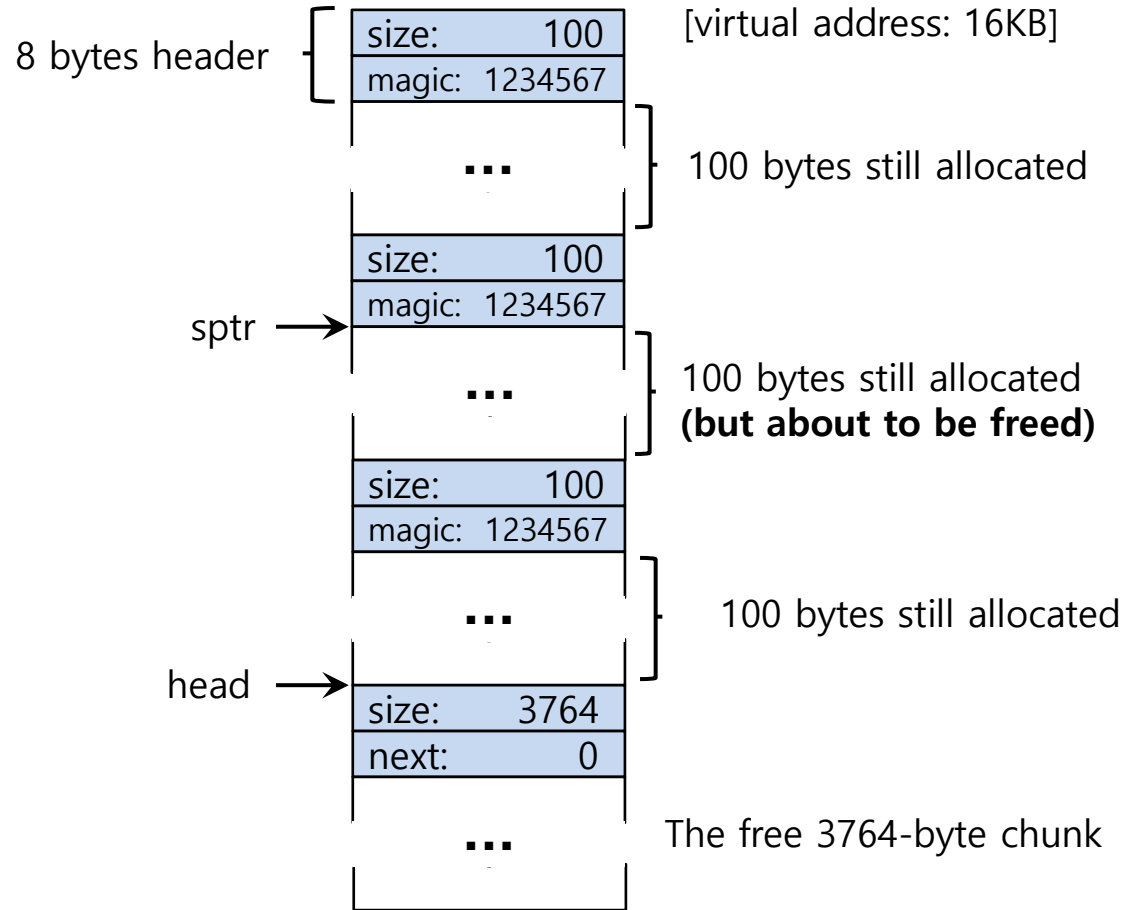
$$4088 \text{ minus } 108 = 3980$$

A Heap : After One Allocation



# Free Space With Chunks Allocated

- Suppose 3 allocations of size 100 bytes each happen
- Then, the middle chunk pointed to by `sptr` is freed

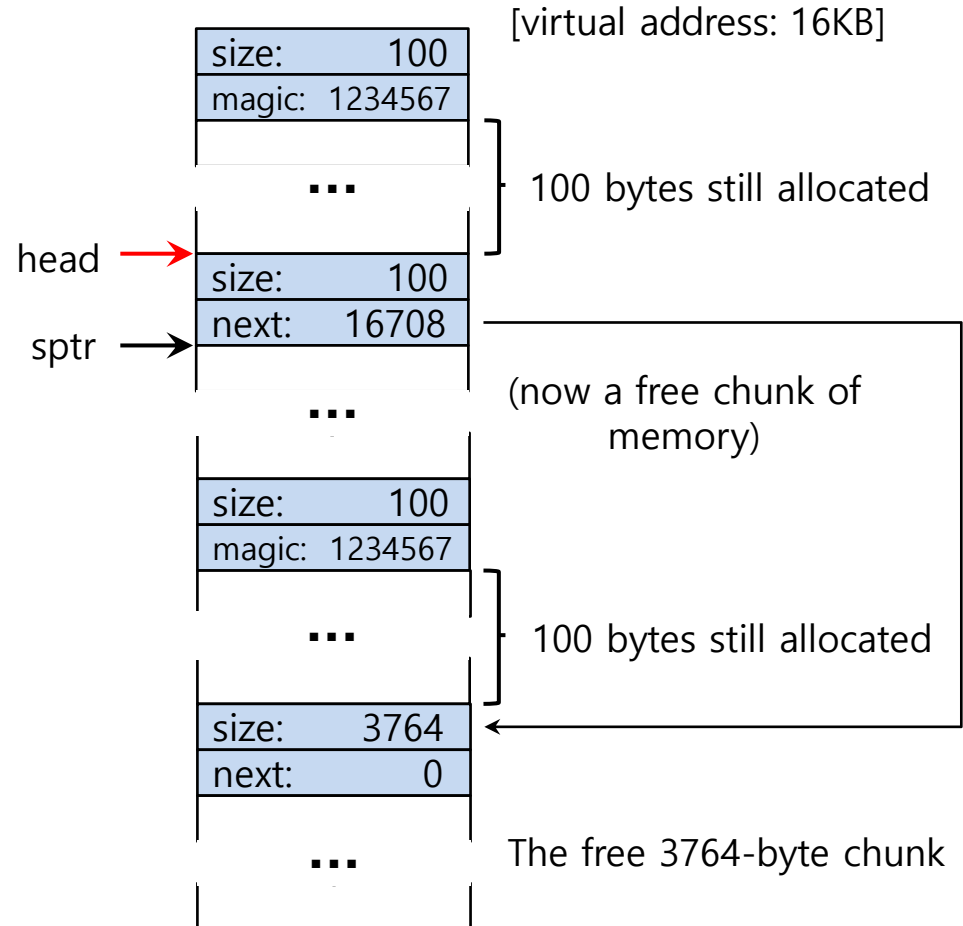


Free Space With Three Chunks Allocated

# Free Space With `free()`

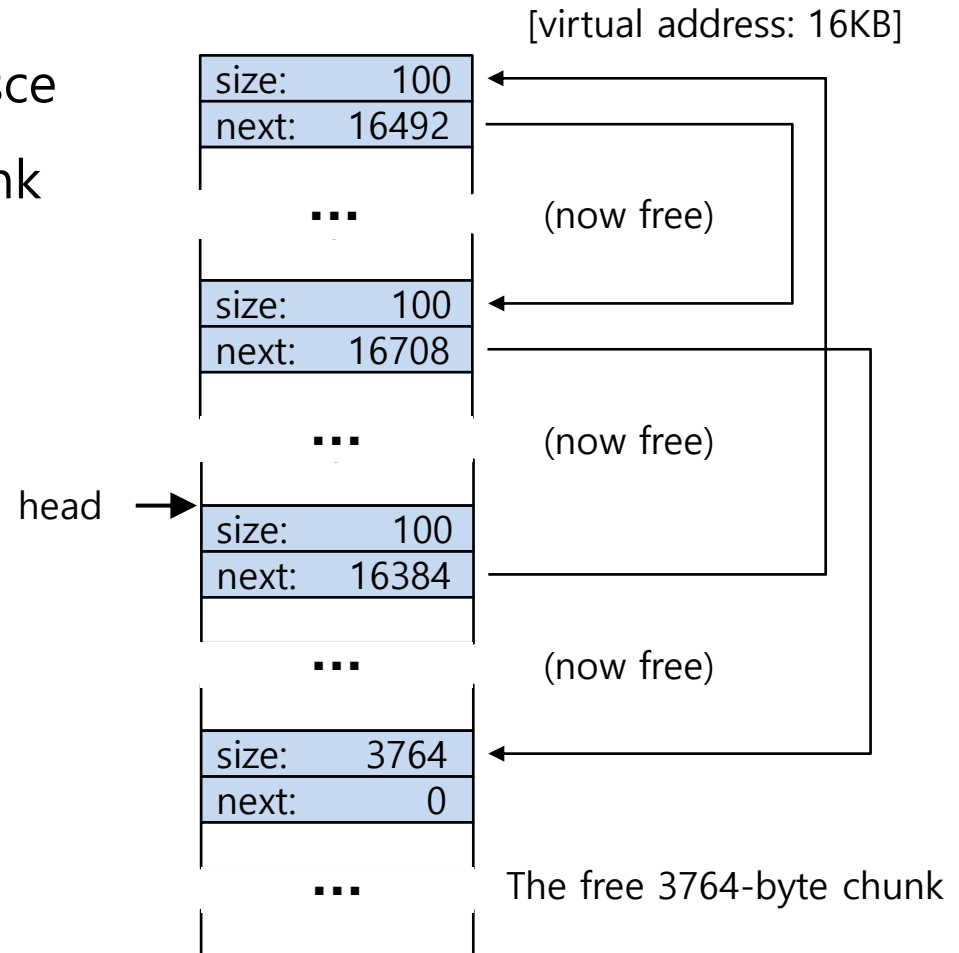
## ▣ Example: `free(sptr)`

- ◆ The 100 bytes chunks is **back into** the free list.
- ◆ The free list header will point the small chunk
- ◆ It now has two non- contiguous elements
- ◆ Free space may be scattered around due to **fragmentation**
- ◆ Cannot satisfy a request for 3800 bytes even though we have the free space



# Free Space With Freed Chunks

- Let's assume that the remaining two in-use chunks are also freed.
- A smart algorithm should coalesce them all into a bigger free chunk




# Managing Free Space: Basic Strategies


- ▣ First Fit: allocate the **first chunk** that is **big enough** for the request
- ▣ Best Fit: allocate free chunk that is closest in size
- ▣ Worst Fit: allocate free chunk that is farthest in size
- ▣ In all these strategies, the requested amount is returned while keeping the **remaining chunk** on the free list.

# Examples of Basic Strategies


- Allocation Request Size 15



- Result of Best-fit: allocate the 20-byte chunk



- Result of Worst-fit: : remaining chunk is bigger and more usable



# Other Approaches: Buddy Allocation

## ❑ Binary Buddy Allocation

- ◆ The allocator **divides free space** by two **until a block** that is big enough to accommodate the request is **found**.
- ◆ E.g., for a request of 7000 bytes, allocate 8 KB chunk

## ❑ Pros: Easy coalescing - if 8KB block and its “buddy” (adjacent block) are free, they can form a 16KB chunk

## ❑ Cons: **internal fragmentation**.

